

# SQL DESIGN AND IMPLEMENTATION

## CONTENT SOURCES:

- ELAMSARI AND NAVATHE, FUNDAMENTALS OF DATABASE MANAGEMENT SYSTEMS
- BRAD LLOYD & MICHELLE ZUKOWSKI'S SLIDES
- Silberschatz–Korth–Sudarshan • Database System Concepts, Fourth Edition

# An Overview of SQL

•2

- Structured Query Language, which is a computer language for storing, manipulating and retrieving data in relational database.
- SQL is an ANSI (American National Standards Institute) standard but there are many different versions of the SQL language.
- SQL is used to communicate with database.
- SQL are used to perform tasks such as create, alter, drop, insert, update and delete etc.
- SQL can execute queries against a database.
- SQL can retrieve data from a database.
- SQL can create new tables, view, index, trigger and cursor etc

# An Overview of SQL

•3

- When a user wants to get some information from a database file, he can issue a **QUERY**.
- A query is a user request to retrieve data or information with a certain condition.
- SQL is a query language that allows user to specify the conditions.
- The user specifies a certain condition.
- The program will go through all the records in the database file and select those records that satisfy (searching) the condition.
- The result of the query will then be stored in form of a table.

# SQL is used for:

•4

- ❑ Data Definition Language(DDL)
- ❑ Data Manipulation Language(DML)
- ❑ Data QUERY Language(DQL)
- ❑ Data Control Language(DCL)
- ❑ Data Transaction Control Language(TCL)

# OVERVIEW OF SQL

•5

## Data Definition Language (DDL)

DDL is the part of SQL that allows a database user to create and restructure database object.

- **CREATE** : Creates a database, new table, a view of a table, index, cursor and trigger or other object in database.
- **ALTER** : Modifies an existing database object, such as a table, view and index etc.
- **DROP** : Deletes an entire table, a view of a table or other object in the database.
- **RENAME**: Used to renaming table, view etc
- **TRUNCATE**: Used to delete record permanente of a table.

# OVERVIEW OF SQL

•6

## Data Manipulation Language (DML)

DML is the part of SQL used to manipulate data within objects of a relational database.

- **INSERT:** used to insert a record on the table.
- **UPDATE:** modification of records in the table.
- **DELETE:** used to remove records in the table.

## □ Data Query Language (DQL):

A query language is a language in which a user requests information from a table.

- **SELECT:** used to retrieve certain records from one or more table.

# OVERVIEW OF SQL

•7

## Data Control Language (DCL):

SQL allow you to control access to data within the database. DCL normally used to create objects related to user access and also control the distribution of privileges among users.

- **GRANT:** is used to grant both system-level and object-level privileges to an existing database user accounts.
- **REVOKE:** To removes privileges that have been granted to database users.

# OVERVIEW OF SQL

•8

## Transaction Control Language (TCL):

commands are used to manage transactions in the database.

These are used to manage the changes made by DML Statement.

- **COMMIT:** Used to permanently save any database transactions.
- **ROLLBACK:** Used restores the database to last committed state.
- **SAVEPOINT:** Used to temporarily save a transaction.



# SQL is a Relational Database

•9

- Represent all info in database as tables
- Keep logical representation of data independent from its physical storage characteristics
- Use one high-level language for structuring, querying, and changing info in the database
- Support the main relational operations
- Support alternate ways of looking at data in tables
- Provide a method for differentiating between unknown values and nulls (zero or blank)
- Support Mechanisms for integrity, authorization, transactions, and recovery

# What Can SQL do?

•10

- ❑ SQL can execute queries against a database
- ❑ SQL can retrieve data from a database
- ❑ SQL can insert records in a database
- ❑ SQL can update records in a database
- ❑ SQL can delete records from a database
- ❑ SQL can create new databases
- ❑ SQL can create new tables in a database
- ❑ SQL can create stored procedures in a database
- ❑ SQL can create views in a database
- ❑ SQL can set permissions on tables, procedures, and views

# SQL DML and DDL

•11

- ❑ SQL can be divided into two parts: The Data Manipulation Language (DML) and the Data Definition Language (DDL).
- ❑ The query and update commands form the DML part of SQL:

SELECT - extracts data from a database

UPDATE - updates data in a database

DELETE - deletes data from a database

INSERT INTO - inserts new data into a database

# The most important DDL statements in SQL are:

•12

- CREATE DATABASE - creates a new database
- ALTER DATABASE - modifies a database
- CREATE TABLE - creates a new table
- ALTER TABLE - modifies a table
- DROP TABLE - deletes a table
- CREATE INDEX - creates an index (search key)
- DROP INDEX - deletes an index

# SQL CREATE TABLE Syntax

•13

```
CREATE TABLE table_name  
(  
  column_name1 data_type constraint if any,  
  column_name2 data_type constraint if any,  
  column_name3 data_type constraint if any,  
  ....  
);
```

```
CREATE TABLE Persons  
(  
  P_Id int,  
  LastName varchar(255),  
  FirstName varchar(255),  
  Address varchar(255),  
  City varchar(255)  
)
```

# Specifying SQL Constraints

•14

- Constraints are used to limit the type of data that can go into table.
- NOT NULL
- UNIQUE
- PRIMARY KEY
- FOREIGN KEY
- CHECK
- DEFAULT

# NOT NULL Constraints in SQL

•15

**NOT NULL** : By default, a column can hold NULL.  
If you not want to allow NULL value in a column,  
you will want to place a constraint on this column.

```
CREATE TABLE Customer  
(  
  SID integer NOT NULL,  
  Last_Name varchar (30) NOT NULL,  
  First_Name varchar(30)  
);
```

# UNIQUE

•16

- The UNIQUE constraint ensures that all values in a column are distinct.

```
CREATE TABLE Customer
```

```
(
```

```
SID integer Unique,
```

```
Last_Name varchar (30),
```

```
First_Name varchar(30)
```

```
);
```



# SQL UNIQUE Constraint

•17

- ❑ The UNIQUE and PRIMARY KEY constraints both provide a guarantee for uniqueness for a column or set of columns.
- ❑ A PRIMARY KEY constraint automatically has a UNIQUE constraint defined on it.
- ❑ You can have many UNIQUE constraints per table, but only one PRIMARY KEY constraint per table.

# CHECK

•18

- The CHECK constraint ensures that all values in a column satisfy certain conditions.

```
CREATE TABLE Customer  
(  
  SID integer CHECK (SID > 0),  
  Last_Name varchar (30),  
  First_Name varchar(30)  
);
```

# Creating Primary key constraint

•19

```
CREATE TABLE Persons (  
  P_Id int NOT NULL,  
  LastName varchar(255) NOT NULL,  
  FirstName varchar(255),  
  Address varchar(255),  
  City varchar(255),  
  CONSTRAINT pk_PersonID PRIMARY KEY (P_Id,LastName)  
)
```

# To DROP a PRIMARY KEY Constraint

•20

```
ALTER TABLE Persons DROP PRIMARY KEY
```

# Creating Foreign Key

•21

- A foreign key is a field (or fields) that points to the primary key of another table.
- The purpose of the foreign key is to ensure referential integrity of the data.

```
CREATE TABLE ORDERS
```

```
(  
  Order_ID integer primary key,  
  Order_Date date,  
  Customer_SID integer references CUSTOMER(SID)  
);
```

# Naming the constraints

•22

```
CREATE TABLE Persons
(
  P_Id int NOT NULL,
  LastName varchar(255) NOT NULL,
  FirstName varchar(255),
  Address varchar(255),
  City varchar(255),
  CONSTRAINT uc_PersonID UNIQUE (P_Id,LastName)
)
```

# Using Alter

•23

## Altering Table

```
ALTER TABLE Persons  
ADD PRIMARY KEY (P_Id)
```

---

## Altering Constraints

```
ALTER TABLE Persons  
ADD CONSTRAINT pk_PersonID PRIMARY KEY  
(P_Id,LastName)
```

# SELECT... FROM... WHERE...

•24

SELECT <attribute list>

FROM <table list>

WHERE <condition>

- <attribute list> is a list of attribute names
- <table list> is a list of the relation names
- <condition> is a conditional (Boolean) expression



## EMPLOYEE

FNAME	MINIT	LNAME	<u>SSN</u>	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
-------	-------	-------	------------	-------	---------	-----	--------	----------	-----

## DEPARTMENT

DNAME	<u>DNUMBER</u>	MGRSSN	MGRSTARTDATE
-------	----------------	--------	--------------

## DEPT\_LOCATIONS

<u>DNUMBER</u>	<u>DLOCATION</u>
----------------	------------------

## PROJECT

PNAME	<u>PNUMBER</u>	PLOCATION	DNUM
-------	----------------	-----------	------

## WORKS\_ON

<u>ESSN</u>	<u>PNO</u>	HOURS
-------------	------------	-------

## DEPENDENT

<u>ESSN</u>	<u>DEPENDENT_NAME</u>	SEX	BDATE	RELATIONSHIP
-------------	-----------------------	-----	-------	--------------

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1

					DEPT_LOCATIONS	DNUMBER	DLOCATION
DEPARTMENT						1	Houston
						4	Stafford
						5	Bellaire
						5	Sugarland
						5	Houston
		DNAME	DNUMBER	MGRSSN	MGRSTARTDATE		
		Research	5	333445555	1988-05-22		
		Administration	4	987654321	1995-01-01		
		Headquarters	1	888665555	1981-06-19		

WORKS_ON	ESSN	PNO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40.0
	453453453	1	20.0
	453453453	2	20.0
	333445555	2	10.0
	333445555	3	10.0
	333445555	10	10.0
	333445555	20	10.0
	999887777	30	30.0
	999887777	10	10.0
	987987987	10	35.0
	987987987	30	5.0
	987654321	30	20.0
	987654321	20	15.0
	888665555	20	null

PROJECT	PNAME	PNUMBER	PLOCATION	DNUM
	ProductX	1	Bellaire	5
	ProductY	2	Sugarland	5
	ProductZ	3	Houston	5
	Computerization	10	Stafford	4
	Reorganization	20	Houston	1
	Newbenefits	30	Stafford	4

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATIONSHIP
	333445555	Alice	F	1986-04-05	DAUGHTER
	333445555	Theodore	M	1983-10-25	SON
	333445555	Joy	F	1958-05-03	SPOUSE
	987654321	Abner	M	1942-02-28	SPOUSE
	123456789	Michael	M	1988-01-04	SON
	123456789	Alice	F	1988-12-30	DAUGHTER
	123456789	Elizabeth	F	1967-05-05	SPOUSE

# Simple SQL Queries

- Query 0: Retrieve the birthdate and address of the employee whose name is 'John B. Smith'.

```
Q0:  SELECT BDATE, ADDRESS
      FROM      EMPLOYEE
      WHERE FNAME='John' AND MINIT='B'
      AND      LNAME='Smith'
```

(a)	<u>Bdate</u>	<u>Address</u>
	1965-01-09	731 Fondren, Houston, TX

# Simple SQL Queries (cont.)

- Query 1: Retrieve the name and address of all employees who work for the 'Research' department.

```
Q1: SELECT  FNAME, LNAME, ADDRESS
      FROM    EMPLOYEE, DEPARTMENT
      WHERE   DNAME='Research' AND DNUMBER=DNO
```

- (DNAME='Research') is a *selection condition* (corresponds to a SELECT operation in relational algebra)
- (DNUMBER=DNO) is a *join condition* (corresponds to a JOIN operation in relational algebra)

(b)

<u>Fname</u>	<u>Lname</u>	<u>Address</u>
John	Smith	731 Fondren, Houston, TX
Franklin	Wong	638 Voss, Houston, TX
Ramesh	Narayan	975 Fire Oak, Humble, TX
Joyce	English	5631 Rice, Houston, TX

# Simple SQL Queries (cont.)

Query 2: For every project located in 'Stafford', list the project number, the controlling department number, and the department manager's last name, address, and birthdate.

```
Q2: SELECT  PNUMBER, DNUM, LNAME, BDATE, ADDRESS
      FROM    PROJECT, DEPARTMENT, EMPLOYEE
      WHERE   DNUM=DNUMBER AND MGRSSN=SSN          AND
             PLOCATION='Stafford'
```

- In Q2, there are *two* join conditions
- The join condition DNUM=DNUMBER relates a project to its controlling department
- The join condition MGRSSN=SSN relates the controlling department to the employee who manages that department

c)

<u>Pnumber</u>	<u>Dnum</u>	<u>Lname</u>	<u>Address</u>	<u>Bdate</u>
10	4	Wallace	291 Berry, Bellaire, TX	1941-06-20
30	4	Wallace	291 Berry, Bellaire, TX	1941-06-20

EMPLOYEE	FNAME	MINIT	LNAME	SSN	BDATE	ADDRESS	SEX	SALARY	SUPERSSN	DNO
	John	B	Smith	123456789	1965-01-09	731 Fondren, Houston, TX	M	30000	333445555	5
	Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
	Alicia	J	Zelaya	999887777	1968-07-19	3321 Castle, Spring, TX	F	25000	987654321	4
	Jennifer	S	Wallace	987654321	1941-06-20	291 Berry, Bellaire, TX	F	43000	888665555	4
	Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
	Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5
	Ahmad	V	Jabbar	987987987	1969-03-29	980 Dallas, Houston, TX	M	25000	987654321	4
	James	E	Borg	888665555	1937-11-10	450 Stone, Houston, TX	M	55000	null	1

					DEPT_LOCATIONS	DNUMBER	DLOCATION
						1	Houston
						4	Stafford
						5	Bellaire
						5	Sugarland
						5	Houston
DEPARTMENT	DNAME	DNUMBER	MGRSSN	MGRSTARTDATE			
	Research	5	333445555	1988-05-22			
	Administration	4	987654321	1995-01-01			
	Headquarters	1	888665555	1981-06-19			

WORKS_ON	ESSN	PNO	HOURS
	123456789	1	32.5
	123456789	2	7.5
	666884444	3	40.0
	453453453	1	20.0
	453453453	2	20.0
	333445555	2	10.0
	333445555	3	10.0
	333445555	10	10.0
	333445555	20	10.0
	999887777	30	30.0
	999887777	10	10.0
	987987987	10	35.0
	987987987	30	5.0
	987654321	30	20.0
	987654321	20	15.0
	888665555	20	null

PROJECT	PNAME	PNUMBER	PLOCATION	DNUM
	ProductX	1	Bellaire	5
	ProductY	2	Sugarland	5
	ProductZ	3	Houston	5
	Computerization	10	Stafford	4
	Reorganization	20	Houston	1
	Newbenefits	30	Stafford	4

DEPENDENT	ESSN	DEPENDENT_NAME	SEX	BDATE	RELATIONSHIP
	333445555	Alice	F	1986-04-05	DAUGHTER
	333445555	Theodore	M	1983-10-25	SON
	333445555	Joy	F	1958-05-03	SPOUSE
	987654321	Abner	M	1942-02-28	SPOUSE
	123456789	Michael	M	1988-01-04	SON
	123456789	Alice	F	1988-12-30	DAUGHTER
	123456789	Elizabeth	F	1967-05-05	SPOUSE

# USE OF \*

- To retrieve all the attribute values of the selected tuples, a \* is used, which stands for *all the attributes*

Examples:

```
Q1C:  SELECT  *  
      FROM EMPLOYEE  
      WHERE   DNO=5
```

(g)

<u>Fname</u>	<u>Minit</u>	<u>Lname</u>	<u>Ssn</u>	<u>Bdate</u>	<u>Address</u>	<u>Sex</u>	<u>Salary</u>	<u>Super_ssn</u>	<u>Dno</u>
John	B	Smith	123456789	1965-09-01	731 Fondren, Houston, TX	M	30000	333445555	5
Franklin	T	Wong	333445555	1955-12-08	638 Voss, Houston, TX	M	40000	888665555	5
Ramesh	K	Narayan	666884444	1962-09-15	975 Fire Oak, Humble, TX	M	38000	333445555	5
Joyce	A	English	453453453	1972-07-31	5631 Rice, Houston, TX	F	25000	333445555	5

# DISTINCT


- SQL does not treat a relation as a set; *duplicate tuples can appear*
- To eliminate duplicate tuples in a query result, the keyword **DISTINCT** is used

Q11: SELECT SALARY  
FROM EMPLOYEE

Q11A: SELECT DISTINCT SALARY  
FROM EMPLOYEE

(a)

SALARY

30000  
40000  
25000  
43000  
38000  
  
55000

(b)

SALARY

30000  
40000  
25000  
43000  
38000  
55000



# GROUPING



To apply the aggregate functions *to subgroups of tuples in a relation*

Each subgroup of tuples consists of the set of tuples that have *the same value* for the *grouping attribute(s)*

The function is applied to each subgroup independently

# GROUPING (cont.)

- Query 20: **For each department**, retrieve the department number, the number of employees in the department, and their average salary.

**Q20: SELECT    DNO, COUNT (\*), AVG (SALARY)  
FROM EMPLOYEE  
GROUP BY DNO**

- In Q20, the EMPLOYEE tuples are divided into groups--each group **having the same value for the grouping attribute DNO**
- The COUNT and AVG functions are applied to each such group of tuples **separately**
- The SELECT-clause includes only the grouping attribute and the functions to be applied on each group of tuples
- A join condition can be used in conjunction with grouping

# GROUPING (cont.)

- Query 21: For each project, retrieve the project number, project name, and the number of employees who work on that project

- ```
Q21:  SELECT  PNUMBER, PNAME, COUNT (*)  
        FROM PROJECT, WORKS_ON  
        WHERE  PNUMBER=PNO  
        GROUP BY  PNUMBER, PNAME
```

- In this case, the grouping and functions are applied *after* the joining of the two relations

# SUBSTRING COMPARISON

- The **LIKE** comparison operator is used to compare **partial strings**
- Two reserved characters are used:

|                 |                                           |
|-----------------|-------------------------------------------|
| 1. '% ' or '* ' | replace an arbitrary number of characters |
| 2. '_ '         | replaces a single arbitrary character     |

# SUBSTRING COMPARISON

## (cont.)

Query 25: Retrieve all employees whose address is in Houston, Texas.  
Here, the value of the ADDRESS attribute must contain the substring 'Houston,TX'.

```
Q25:  SELECT  FNAME, LNAME
        FROM    EMPLOYEE
        WHERE     ADDRESS LIKE '%Houston,TX%'
```

# SUBSTRING COMPARISON (cont.)

- Query 26: Retrieve all employees who were born during the 1950s.

```
Q26:  SELECT  FNAME, LNAME  
      FROM    EMPLOYEE  
      WHERE   BDATE LIKE '____5_'
```

The LIKE operator allows us to get around the fact that each value is considered atomic and indivisible; hence, in SQL, character string attribute values are not atomic

Here, '5' must be the 8th character of the string (according to our format for date), so the BDATE value is '\_\_\_\_5\_', with each underscore as a place holder for a single arbitrary character.

# ARITHMETIC OPERATIONS

- The standard arithmetic operators '+', '-', '\*', and '/' can be applied to numeric values in an SQL query result
- Query 27: Show the effect of giving all employees who work on the 'ProductX' project a 10% raise.

```
Q27:  SELECT      FNAME, LNAME, 1.1*SALARY
        FROM      EMPLOYEE, WORKS_ON, PROJECT
        WHERE      SSN=ESSN AND PNO=PNUMBER AND
PNAME='ProductX'
```

# ORDER BY

- The ORDER BY clause is used to sort the tuples in a query result **based on the values of some attribute(s)**
- Query 28: Retrieve a list of employees and the projects each works in, ordered by the employee's department, and within each department ordered alphabetically by employee last name.

```
Q28: SELECT      DNAME, LNAME, FNAME, PNAME
      FROM        DEPARTMENT, EMPLOYEE, WORKS_ON, PROJECT
      WHERE       DNUMBER=DNO AND SSN=ESSN AND PNO=PNUMBER
      ORDER BY    DNAME, LNAME
```



# ORDER BY (cont.)

- The default order is in **ascending** order of values
- We can specify the keyword **DESC** if we want a descending order;
- The keyword **ASC** can be used to explicitly specify ascending order, even though it is the default

# NULLS IN SQL QUERIES

- SQL allows queries that check if a value is NULL
  - (missing / undefined /not applicable)
- SQL uses IS or IS NOT to compare NULLs because it considers each NULL value distinct from other NULL values, so equality comparison is not appropriate .
- Query 14: Retrieve the names of all employees who do not have supervisors.

```
□ Q14: SELECT  FNAME, LNAME
           FROM    EMPLOYEE
           WHERE   SUPERSSN IS NULL
```

Note: If a join condition is specified, tuples with NULL values for the join attributes are not included in the result

# Summary of SQL Queries

- A query in SQL can consist of up to six clauses, but only the first two, SELECT and FROM, are mandatory. The clauses are specified in the following order:
- **SELECT <attribute list>**  
**FROM <table list>**  
**[WHERE <condition>]**  
**[GROUP BY <grouping attribute(s)>]**  
**[HAVING <group condition>]**  
**[ORDER BY <attribute list>]**

# Summary of SQL Queries (cont.)

- The SELECT-clause lists the attributes or functions to be retrieved
- The FROM-clause specifies all relations (or aliases) needed in the query but not those needed in nested queries
- The WHERE-clause specifies the conditions for selection and join of tuples from the relations specified in the FROM-clause
- GROUP BY specifies grouping attributes
- HAVING specifies a condition for selection of groups
- ORDER BY specifies an order for displaying the result of a query
- A query is evaluated by first applying the WHERE-clause, then GROUP BY and HAVING, and finally the SELECT-clause

# Specifying Updates in SQL

- There are three SQL commands to modify the database; INSERT, DELETE, and UPDATE

# INSERT

- Used to add one or more tuples to a relation
- Attribute values should be listed in the same order as the attributes were specified in the CREATE TABLE command

# INSERT (cont.)

- Example:

```
U1: INSERT INTO EMPLOYEE VALUES ('Richard','K','Marini',  
'653298653', '30-DEC-52','98 Oak Forest,Katy,TX', 'M', 7000,'987654321', 4 )
```

- An alternate form of INSERT specifies explicitly the attribute names that correspond to the values in the new tuple
- Example: Insert a tuple for a new EMPLOYEE for **whom we only know** the FNAME, LNAME, and SSN attributes.

```
U1A: INSERT INTO EMPLOYEE (FNAME, LNAME, SSN)  
VALUES ('Richard', 'Marini', '653298653')
```

# INSERT (cont.)

- Note: Only the constraints specified in the DDL commands are automatically enforced by the DBMS when updates are applied to the database
- Another variation of INSERT allows insertion of *multiple tuples* resulting from a query into a relation



# INSERT (cont.)

- Example: Suppose we want to create a temporary table that has the name, number of employees, and total salaries for each department. A table DEPTS\_INFO is created by U3A, and is loaded with the summary information retrieved from the database by the query in U3B.

```
U3A:  CREATE TABLE DEPTS_INFO
      (DEPT_NAME    VARCHAR(10),
       NO_OF_EMPS   INTEGER,
       TOTAL_SAL    INTEGER);
```

```
U3B:  INSERT INTO   DEPTS_INFO (DEPT_NAME,
NO_OF_EMPS, TOTAL_SAL)
      SELECT        DNAME, COUNT (*), SUM (SALARY)
      FROM          DEPARTMENT, EMPLOYEE
      WHERE         DNUMBER=DNO
      GROUP BY      DNAME ;
```

# INSERT (cont.)

- Note: The DEPTS\_INFO table may not be up-to-date if we change the tuples in either the DEPARTMENT or the EMPLOYEE relations *after* issuing U3B. We have to create a view (see later) to keep such a table up to date.

# DELETE

- ❑ Removes tuples from a relation
- ❑ Includes a WHERE-clause to select the tuples to be deleted
- ❑ Tuples are deleted from only *one table* at a time (unless CASCADE is specified on a referential integrity constraint)
- ❑ A missing WHERE-clause specifies that *all tuples* in the relation are to be deleted; the table then becomes an empty table
- ❑ The number of tuples deleted depends on the number of tuples in the relation that satisfy the WHERE-clause
- ❑ Referential integrity should be enforced

# DELETE (cont.)

U4A: DELETE FROM EMPLOYEE  
WHERE LNAME='Brown'

U4B: DELETE FROM EMPLOYEE  
WHERE SSN='123456789'

U4C: DELETE FROM EMPLOYEE  
WHERE DNO IN (SELECT  
DNUMBER  
FROM DEPARTMENT  
WHERE DNAME='Research')

U4D: DELETE FROM EMPLOYEE

# UPDATE

- Used to modify attribute values of one or more selected tuples
- A WHERE-clause selects the tuples to be modified
- An additional SET-clause specifies the attributes to be modified and their new values
- Each command modifies tuples *in the same relation*
- Referential integrity should be enforced

# UPDATE (cont.)

- Example: Change the location and controlling department number of project number 10 to 'Bellaire' and 5, respectively.

```
U5: UPDATE PROJECT
     SET      PLOCATION = 'Bellaire', DNUM = 5
     WHERE    PNUMBER=10
```

# UPDATE (cont.)

- Example: Give all employees in the 'Research' department a 10% raise in salary.

```
U6: UPDATE      EMPLOYEE
    SET      SALARY = SALARY *1.1
    WHERE      DNO IN (SELECT  DNUMBER
                          FROM      DEPARTMENT
                          WHERE  DNAME='Research')
```

- In this request, the modified SALARY value depends on the original SALARY value in each tuple
- The reference to the SALARY attribute on the right of = refers to the old SALARY value before modification
- The reference to the SALARY attribute on the left of = refers to the new SALARY value after modification

# Thankyou